


最短通路




weighted graph (加权图、有权图)

- ◆ 在某些时候某些场合，并非图的所有边都一样长。出于某些原因和目的，需要给图的每条边加权(某种意义上的值、长度)，也即给边赋一个值，称为边的长度。
- ◆ 举例说明一下加权的必要性。
- ◆ Weights can also be attached to the vertices instead of the edges or can be attached to both vertices and edges. The resulting graph is called a *weighted graph*.

Weighted Graphs



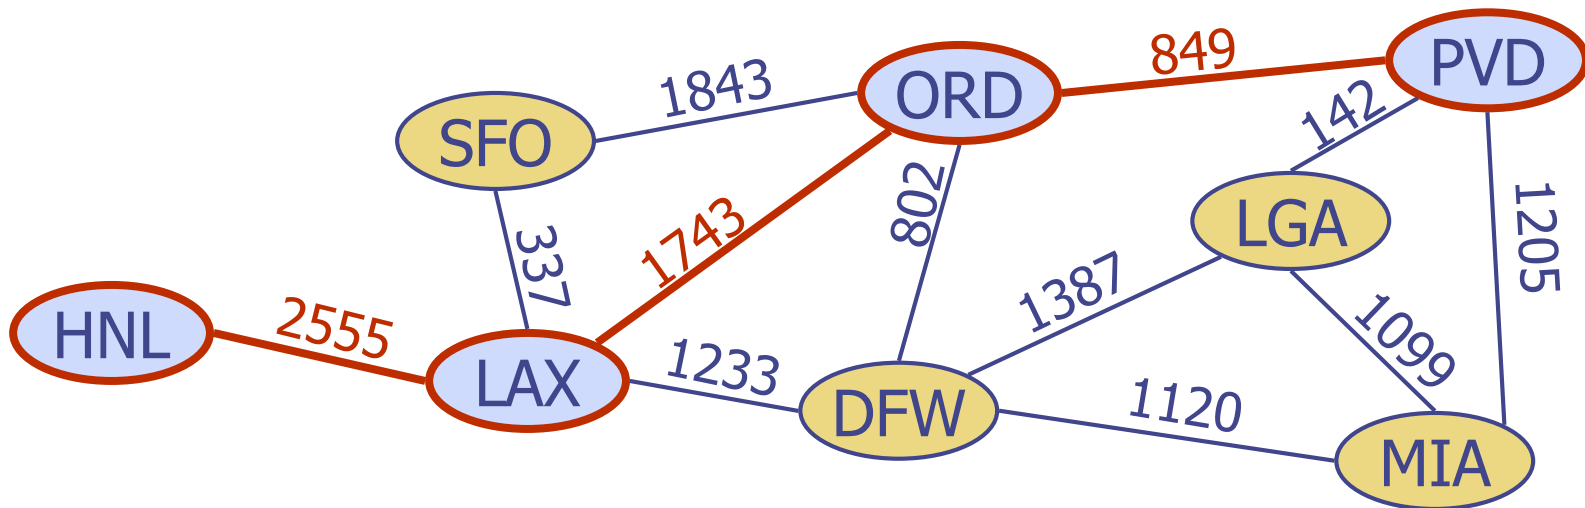
- ◆ 加权图中的每条边被赋予一个数值（权）
- ◆ 权可以是距离，时间，成本，费用，带宽，吞吐量等等。
- ◆ Example:
 - In a flight route graph, the weight of an edge represents the distance in miles between the endpoint airports



Shortest Path Problem



- ◆ 最小通路问题：给定有权图中的两个不同结点，寻找两个点之间总权最小的通路
- ◆ E.G: Shortest path between Providence and Honolulu
- ◆ Applications
 - Internet packet routing
 - Flight reservations
 - Driving directions



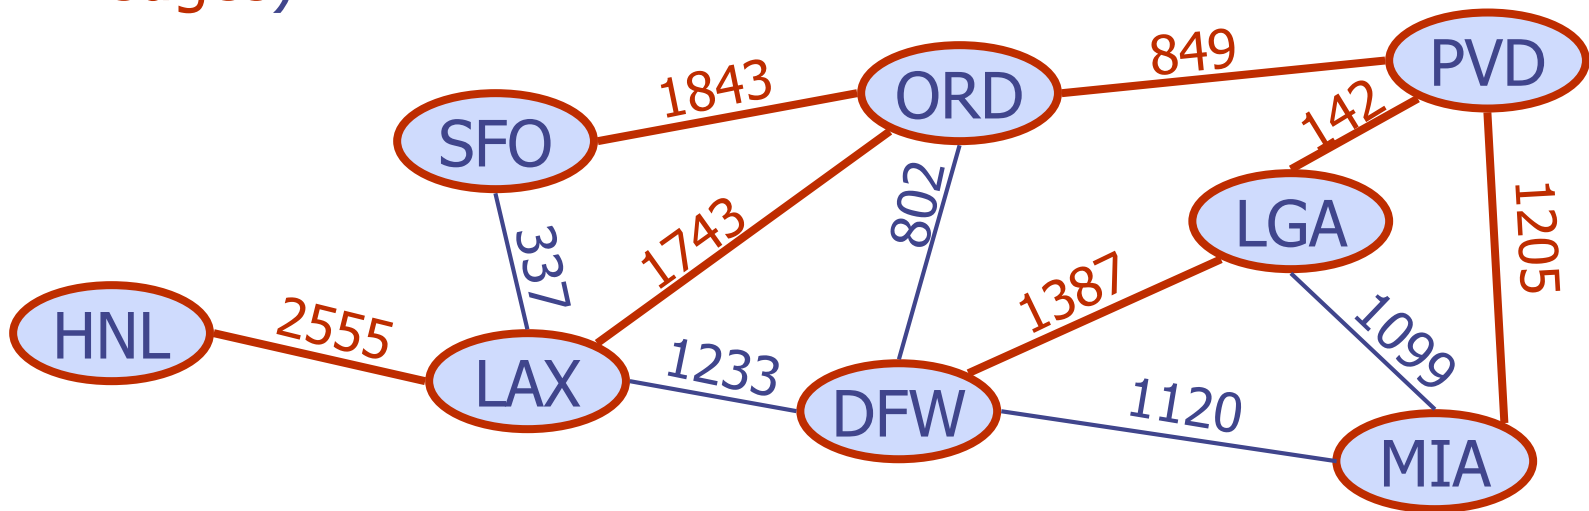
Shortest Path Properties

性质1: 最短通路的子路本身也一定是一条最短通路。
(Why?)

性质2: 连通图中, 存在一颗从一个起始结点到其它所有结点的最短路径的树。

Example:

Tree of shortest paths from Providence (those red edges)



Dijkstra's 算法 (最经典的、最常用的算法)

single-source shortest path problem in graph theory. 图论最短通路问题的**单源算法** (给定一个点到其它所有连接的点的)

Works for both undirected and digraph. 但只适应非负权图。

算法Input: Weighted graph $G=(V,E,f)$ and source vertex $s \in V$, such that all edge weights are nonnegative

算法Output: Lengths of shortest paths (or the shortest paths themselves) **from a given source vertex $s \in V$ to all other vertices**

Dijkstra's Algorithm (迪克斯特拉单源算法)

- ◆ 距离：一个结点到 v 到另一结点 s 的距离是 v,s 间的最短通路的长度，这里的长度是路的所有边的权之和。
- ◆ Dijkstra's algorithm 计算起点 s 到所有其它所有连接的点的距离。
- ◆ Assumptions:
 - 连通
 - 所有权值非负
 - 简单无向图

Dijkstra's Algorithm (迪克斯特拉单源算法)

- ◆ “云”：结点集 V 的子集（从点 s 开始慢慢形成扩张）
- ◆ “云”算法，或者叫“水淹”算法：给每一个结点 v 保存一个临时值 $d(v)$ ，表示在“云”以及云邻接的所有结点形成的子图中从起点 s 到 v 的距离。最终将这“云”扩大到整个图
- ◆ At each step（“云”扩张过程，也即迭代的过程）
 - 开始“云”只包括结点 s 一个点 ($d(s)=0$)
 - 把云外的 $d(u)$ 最小的点 u 加入到云中（也即离“云”最近的点）
 - 更新“云”外与 u 邻接的结点的的标记 $d(v)$.（关键搞清楚如何更新 $d(v)$ ）
 - 当“云”扩张到了整个图，所有的 $d(v)$ 都标注完，任务完成

Edge Relaxation


◆ 第一次给所有与起点 s 邻接的点中离 s 最近的点标注一个距离 $d(v)$ （也即相应的边长）。每一个与 s 邻接的结点标注成相应边的长度；其它所有外围的点标为 ∞

◆ Consider an edge $e = (u, z)$ such that

- u 是最近加入到云中的结点

◆ The relaxation of edge e updates distance $d(z)$ as follows:

$d(z) = \min\{d(z), d(u) + \text{weight}(e)\}$ 逐个更新与云中点 u 邻接的在云外的结点 z 的标注值 $d(z)$ （“云”周边的）




Edge Relaxation

- ◆ 逐个更新与云中点邻接的在云外的结点 z 的标注值 $d(z)$ （“云”周边的）


- ◆ 将“云”周边邻接的点中距离（标注值）最小的点加入到“云”中

- ◆ 直到“云”包含了所有的结点



- ◆ 注：这里所谓的“云”其实是结点集 V 的一个子集。



举例：观察云的扩张过程以及点的值 $d(v)$ 的变化过程




Example (cont.)



Another Dijkstra Animated Example for directed graph (有向图距离)

Initialize:




$Q:$


A	B	C	D	E
0	∞	∞	∞	∞

$S: \{\}$


Dijkstra Animated Example




Dijkstra Animated Example




Dijkstra Animated Example




Dijkstra Animated Example




Dijkstra Animated Example




Dijkstra Animated Example




Dijkstra Animated Example



Dijkstra Animated Example



Dijkstra Animated Example




Dijkstra 算法每次迭代要做的两件事

- ◆ 1. 从子集 S 外（“云”外）的所有与 S 中结点邻接的所有结点中选择标注值 $d(z)$ 最小的结点 u ，加入到 S 中；
 - ◆ 2. 考察对比与 u 结点邻接的在 S 之外的结点的标注值，做可能的修改。
- ◆ 注：一个结点 z 的标注值 $d(z)$ ，当它 $\neq \infty$ 时，所代表的含义是：如果 z 在 S 中，则它是从起点到点 z 的最短路径的长度，也即距离；如果在 S 外则说明有某一条从起点到 z 的路，长度为 $d(z)$. 这个值的不断修改过程就是寻找更短路的过程，直到找到最短的为止。

Why Dijkstra's Algorithm Works

◆ Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.

- Suppose it didn't find all shortest distances. Let F be the first wrong vertex the algorithm processed.
- When the previous node, D , on the true shortest path was considered, its distance was correct.
- But the edge (D, F) was **relaxed** at that time!
- Thus, so long as $d(F) \geq d(D)$ (非负边), F 's distance cannot be wrong. That is, there is no wrong vertex.




Why It Doesn't Work for Negative-Weight Edges

- ◆ Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.

如果一个结点与一条带负权的边关联，被加入到“云”里，就可能导致混乱。

如右图所示：

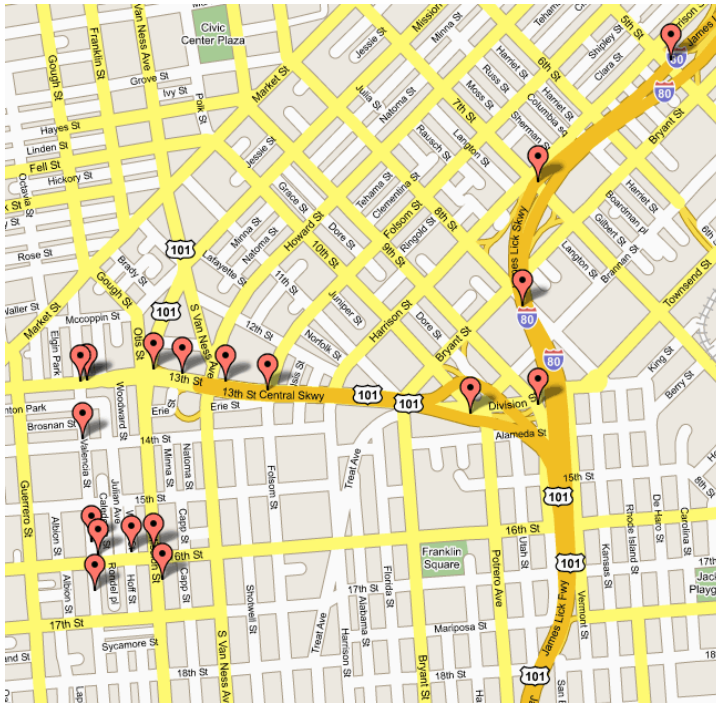


C's true distance is 1, but it is already in the cloud with $d(C)=5$!

Applications of Dijkstra's Algorithm

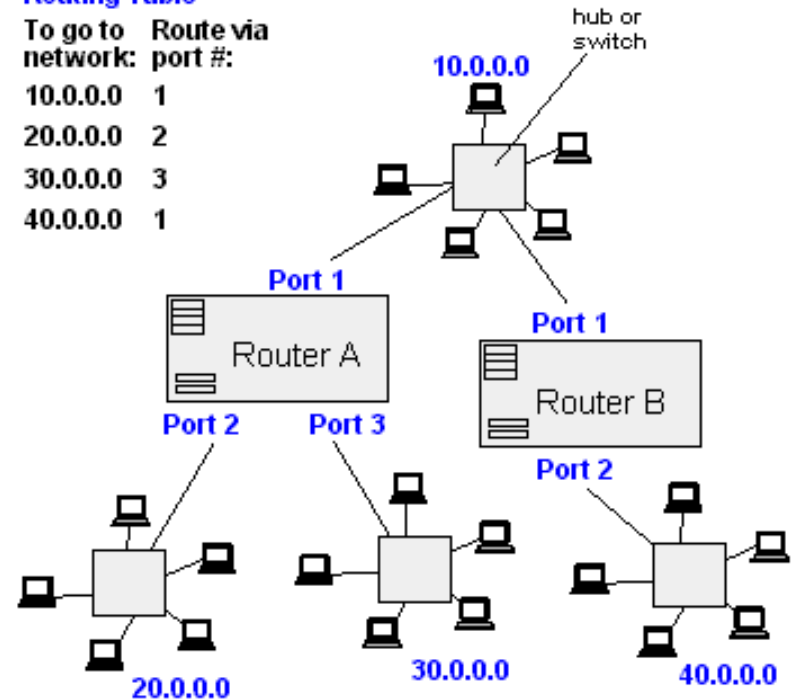
- Traffic Information Systems are most prominent use
- Mapping (Map Quest, Google Maps)
- Routing Systems

From Computer Desktop Encyclopedia
© 1998 The Computer Language Co. Inc.



Router A Routing Table

To go to network:	Route via port #:
10.0.0.0	1
20.0.0.0	2
30.0.0.0	3
40.0.0.0	1





Dijkstra's 算法应用

◆ One particularly relevant: epidemiology

◆ Prof. Lauren Meyers (Biology Dept.) uses networks to model the spread of infectious diseases and design prevention and response strategies. (传染病防控)

◆ Vertices represent individuals, and edges their possible contacts. It is useful to calculate how a particular individual is connected to others.

◆ Knowing the shortest path lengths to other individuals can be a relevant indicator of the potential of a particular individual to infect others.



◆ How to solve the shortest path problem when the graph has negative edges?

◆ Bellman-Ford Algorithm （求含负权图的单源最短路径算法，效率很低，但代码很容易写）


(<http://blog.csdn.net/xu3737284/article/details/8973615>)

◆ DAG-based Algorithm

(http://blog.csdn.net/wall_f/article/details/8204747)

注：这两个算法自己有兴趣的话，上网去搜索学习

Bellman-Ford Algorithm



- ◆ Works even with negative-weight edges
- ◆ Must assume directed edges (有向无环) (for otherwise we would have negative-weight cycles)
- ◆ Iteration i finds all shortest paths that use i edges.
- ◆ Running time: $O(nm)$.
- ◆ Can be extended to detect a negative-weight cycle if it exists
 - How?

Algorithm *BellmanFord*(G, s)

 for all $v \in G.vertices()$

 if $v = s$

$setDistance(v, 0)$

 else

$setDistance(v, \infty)$

 for $i \leftarrow 1$ to $n-1$ do

 for each $e \in G.edges()$

 { relax edge e }

$u \leftarrow G.origin(e)$

$z \leftarrow G.opposite(u, e)$


$r \leftarrow getDistance(u) + weight(e)$

 if $r < getDistance(z)$


$setDistance(z, r)$

Bellman-Ford Example

Nodes are labeled with their $d(v)$ values



DAG-based Algorithm





- ◆ Works even with negative-weight edges
- ◆ Uses topological order
- ◆ Doesn't use any fancy data structures
- ◆ Is much faster than Dijkstra's algorithm
- ◆ Running time: $O(n+m)$.

```
Algorithm DagDistances( $G, s$ )  
  for all  $v \in G.vertices()$   
    if  $v = s$   
      setDistance( $v, 0$ )  
    else  
      setDistance( $v, \infty$ )  
  Perform a topological sort of the vertices  
  for  $u \leftarrow 1$  to  $n$  do {in topological order}  
    for each  $e \in G.outEdges(u)$   
      { relax edge  $e$  }  
       $z \leftarrow G.opposite(u, e)$   
       $r \leftarrow getDistance(u) + weight(e)$   
      if  $r < getDistance(z)$   
        setDistance( $z, r$ )
```

DAG Example

Nodes are labeled with their $d(v)$ values





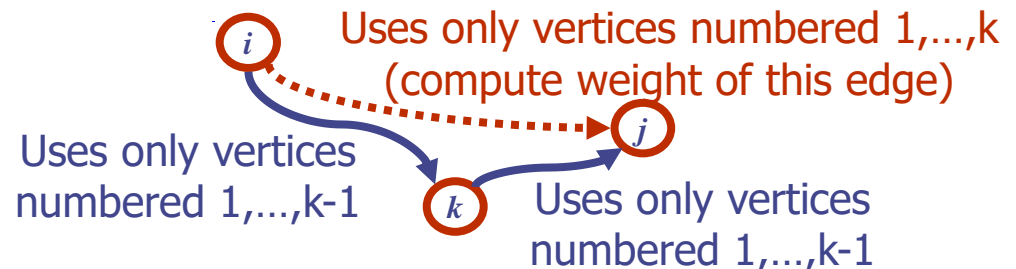
All-Pairs Shortest Paths

- Find the distance between every pair of vertices in a weighted directed graph G .
- We can make n calls to Dijkstra's algorithm (if no negative edges), which takes $O(nm \log n)$ time.
- Likewise, n calls to Bellman-Ford would take $O(n^2m)$ time.
- We can achieve $O(n^3)$ time using dynamic programming (similar to the Floyd-Warshall algorithm).

Algorithm *AllPair*(G) {assumes vertices $1, \dots, n$ }

```

for all vertex pairs  $(i, j)$ 
  if  $i = j$ 
     $D_0[i, i] \leftarrow 0$ 
  else if  $(i, j)$  is an edge in  $G$ 
     $D_0[i, j] \leftarrow \text{weight of edge } (i, j)$ 
  else
     $D_0[i, j] \leftarrow +\infty$ 
for  $k \leftarrow 1$  to  $n$  do
  for  $i \leftarrow 1$  to  $n$  do
    for  $j \leftarrow 1$  to  $n$  do
       $D_k[i, j] \leftarrow \min\{D_{k-1}[i, j], D_{k-1}[i, k] + D_{k-1}[k, j]\}$ 
return  $D_n$ 
  
```



思考Question

- ◆ 对一个普通无向连通图, For a connected simple undirected graph (non-weighted), can you design an algorithm to calculate the distance between a start vertex v_0 to any other vertex v ?
- ◆ *Solution: you can set weight 1 to all edges of the graph*

Traveling Salesman Problem

◆ Introduction to Traveling Salesman Problem

- ◆ 某售货员要到若干城市去推销商品，已知各城市之间的路程(或旅费)。他要选定一条从驻地出发，经过每个城市一次，最后回到驻地的路线，使总的路程(或总旅费)最小。
- ◆ 数学化的问题：在带权完全无向图里，求访问每个顶点一次只一次，且最后返回出发点，总权最小的路。这实质是求完全图里总权最小的哈密尔顿回路。
- ◆ 这是一个NP-问题。

练习

◆ 6.6节 T1, T2